



INSTITUT
DES SCIENCES
DE L'HOMME




Erasmus Mundus
master course in
**Data Mining and
Knowledge Management**
a european master



UNIVERSITÀ DEL PIEMONTE ORIENTALE

UNIVERSITÉ
LUMIÈRE
LYON 2
UNIVERSITÉ DE LYON

MACHINE LEARNING IN PRACTICE



Hussein AL-NATSHEH, CIAPPLE

Jordan Open Source Association

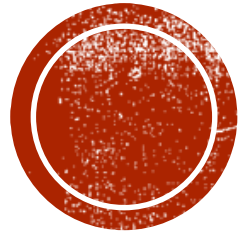
Amman, Aug 22, 2015

Hussein AL-NATSHEH ©2015

AGENDA

- From supervised to unsupervised learning
- Identifying a ML problem?
- Recommended python libraries
- Data preparation
- Case study; Author name disambiguation of CERN digital library (use case)
- References, recommended readings and courses





FROM SUPERVISED TO UNSUPERVISED LEARNING



RECALL SUPERVISED LEARNING

Relation: breast-cancer

No.	age Nominal	menopause Nominal	tumor-size Nominal	inv-nodes Nominal	node-caps Nominal	deg-malig Nominal	breast Nominal	breast-quad Nominal	irradiat Nominal	Class Nominal
1	40-49	premeno	15-19	0-2	yes	3	right	left_up	no	recurrence-events
2	50-59	ge40	15-19	0-2	no	1	right	central	no	no-recurrence-events
3	50-59	ge40	35-39	0-2	no	2	left	left_low	no	recurrence-events
4	40-49	premeno	35-39	0-2	yes	3	right	left_low	yes	no-recurrence-events
5	40-49	premeno	30-34	3-5	yes	2	left	right_up	no	recurrence-events
6	50-59	premeno	25-29	3-5	no	2	right	left_up	yes	no-recurrence-events
7	50-59	ge40	40-44	0-2	no	3	left	left_up	no	no-recurrence-events
8	40-49	premeno	10-14	0-2	no	2	left	left_up	no	no-recurrence-events
9	40-49	premeno	0-4	0-2	no	2	right	right_low	no	no-recurrence-events
10	40-49	ge40	40-44	15-17	yes	2	right	left_up	yes	no-recurrence-events
11	50-59	premeno	25-29	0-2	no	2	left	left_low	no	no-recurrence-events
12	60-69	ge40	15-19	0-2	no	2	right	left_up	no	no-recurrence-events
13	50-59	ge40	30-34	0-2	no	1	right	central	no	no-recurrence-events
14	50-59	ge40	25-29	0-2	no	2	right	left_up	no	no-recurrence-events
15	40-49	premeno	25-29	0-2	no	2	left	left_low	yes	recurrence-events
16	30-39	premeno	20-24	0-2	no	3	left	central	no	no-recurrence-events
17	50-59	premeno	10-14	3-5	no	1	right	left_up	no	no-recurrence-events
18	60-69	ge40	15-19	0-2	no	2	right	left_up	no	no-recurrence-events
19	50-59	premeno	40-44	0-2	no	2	left	left_up	no	no-recurrence-events
20	50-59	ge40	20-24	0-2	no	3	left	left_up	no	no-recurrence-events
21	50-59	lt40	20-24	0-2		1	left	left_low	no	recurrence-events
22	60-69	ge40	40-44	3-5	no	2	right	left_up	yes	no-recurrence-events



UNSUPERVISED LEARNING

Relation: breast-cancer

No.	age Nominal	menopause Nominal	tumor-size Nominal	inv-nodes Nominal	node-caps Nominal	deg-malig Nominal	breast Nominal	breast-quad Nominal	irradiat Nominal	Class Nominal
1	40-49	premeno	15-19	0-2	yes	3	right	left_up	no	recurrence-events
2	50-59	ge40	15-19	0-2	no	1	right	central	no	no-recurrence-events
3	50-59	ge40	35-39	0-2	no	2	left	left_low	no	recurrence-events
4	40-49	premeno	35-39	0-2	yes	3	right	left_low	yes	no-recurrence-events
5	40-49	premeno	30-34	3-5	yes	2	left	right_up	no	recurrence-events
6	50-59	premeno	25-29	3-5	no	2	right	left_up	yes	no-recurrence-events
7	50-59	ge40	40-44	0-2	no	3	left	left_up	no	no-recurrence-events
8	40-49	premeno	10-14	0-2	no	2	left	left_up	no	no-recurrence-events
9	40-49	premeno	0-4	0-2	no	2	right	right_low	no	no-recurrence-events
10	40-49	ge40	40-44	15-17	yes	2	right	left_up	yes	no-recurrence-events
11	50-59	premeno	25-29	0-2	no	2	left	left_low	no	no-recurrence-events
12	60-69	ge40	15-19	0-2	no	2	right	left_up	no	no-recurrence-events
13	50-59	ge40	30-34	0-2	no	1	right	central	no	no-recurrence-events
14	50-59	ge40	25-29	0-2	no	2	right	left_up	no	no-recurrence-events
15	40-49	premeno	25-29	0-2	no	2	left	left_low	yes	recurrence-events
16	30-39	premeno	20-24	0-2	no	3	left	central	no	no-recurrence-events
17	50-59	premeno	10-14	3-5	no	1	right	left_up	no	no-recurrence-events
18	60-69	ge40	15-19	0-2	no	2	right	left_up	no	no-recurrence-events
19	50-59	premeno	40-44	0-2	no	2	left	left_up	no	no-recurrence-events
20	50-59	ge40	20-24	0-2	no	3	left	left_up	no	no-recurrence-events
21	50-59	lt40	20-24	0-2		1	left	left_low	no	recurrence-events
22	60-69	ge40	40-44	3-5	no	2	right	left_up	yes	no-recurrence-events



CLUSTERING

- A flat partition (set of clusters or segments)
- A hierarchical tree or taxonomy (a set of nested partitions)
- Hard or soft (or fuzzy) memberships to clusters

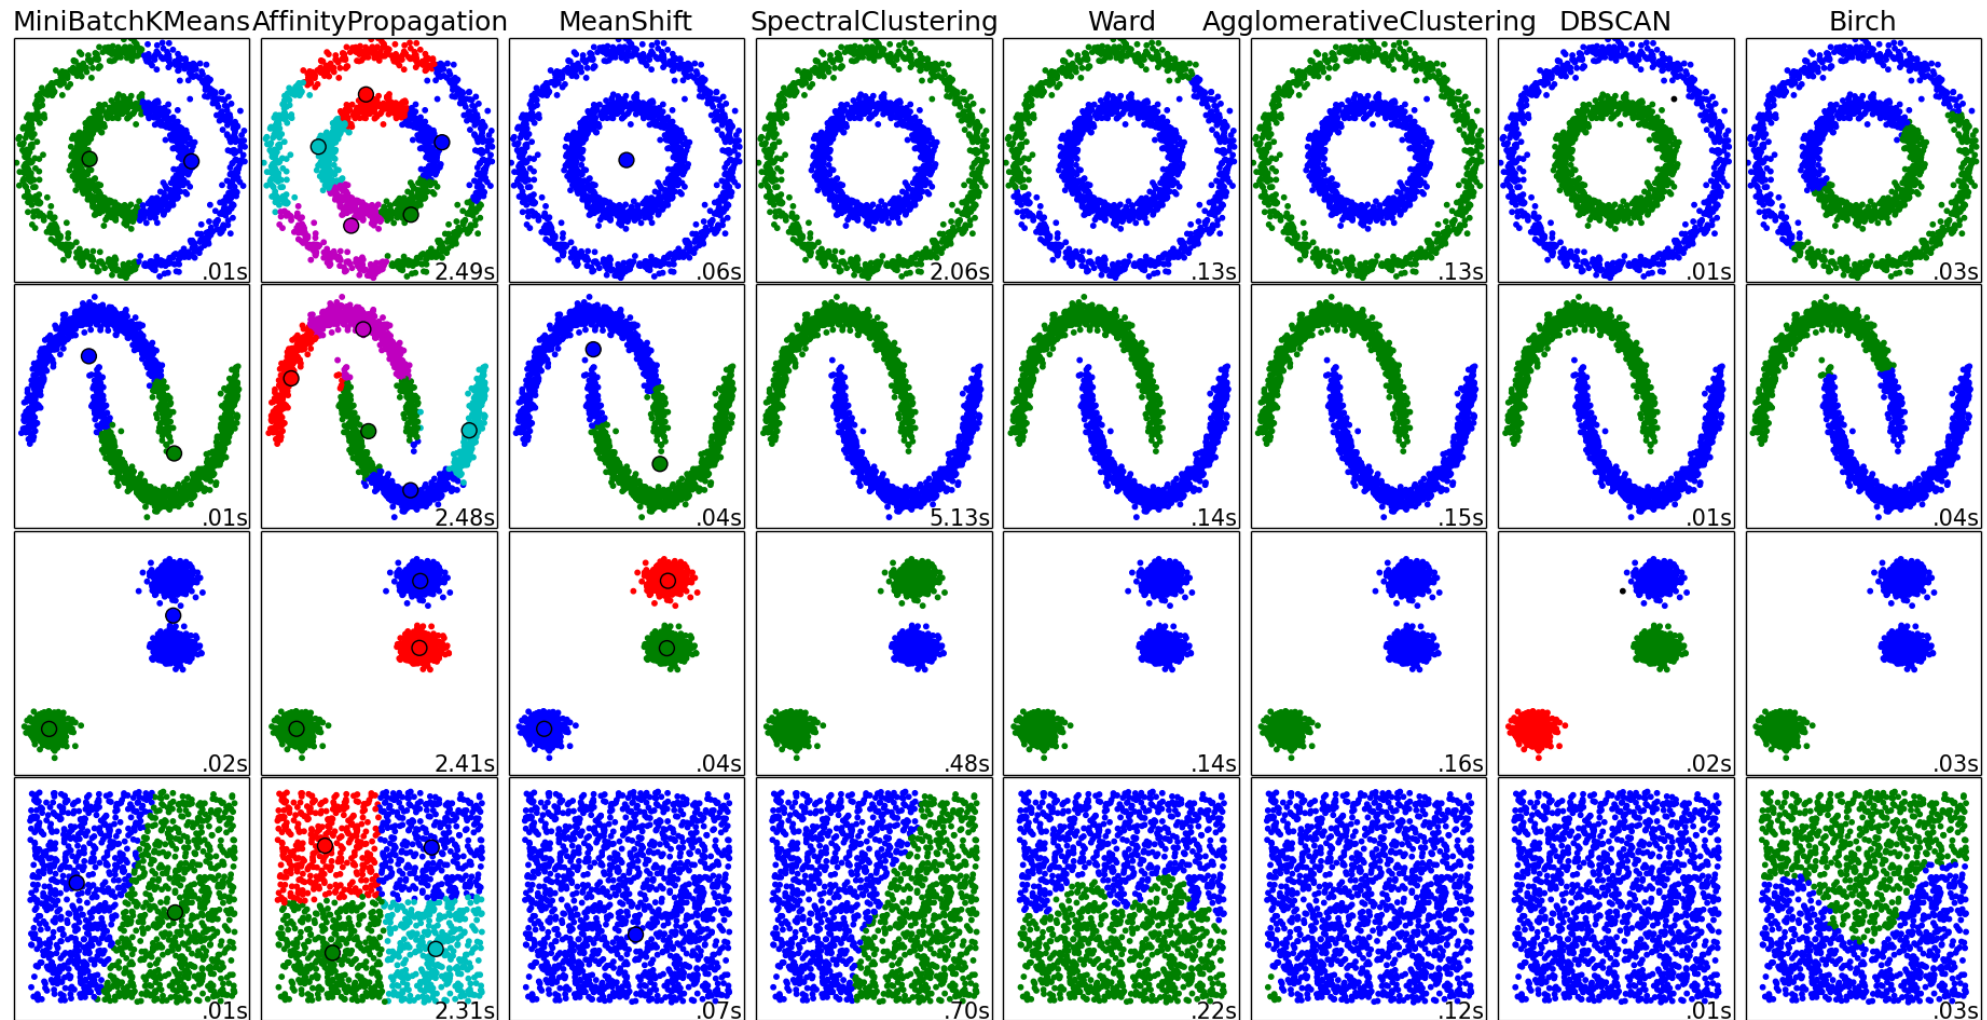


SOME CLUSTERING ALGORITHMS (10F2)

- K-Means, Expectation Maximization
- Hierarchical Agglomerative Clustering (HAC)
- Density-based spatial clustering of applications with noise (DBSCAN)
- Graph-based clustering (Modularity, Weighted communities, Clique percolation)
- Unsupervised Neural Networks (Adaptive Resonance Theory)




SOME CLUSTERING ALGORITHMS (20F2)



Source: <http://scikit-learn.org/stable/modules/clustering.html#clustering>

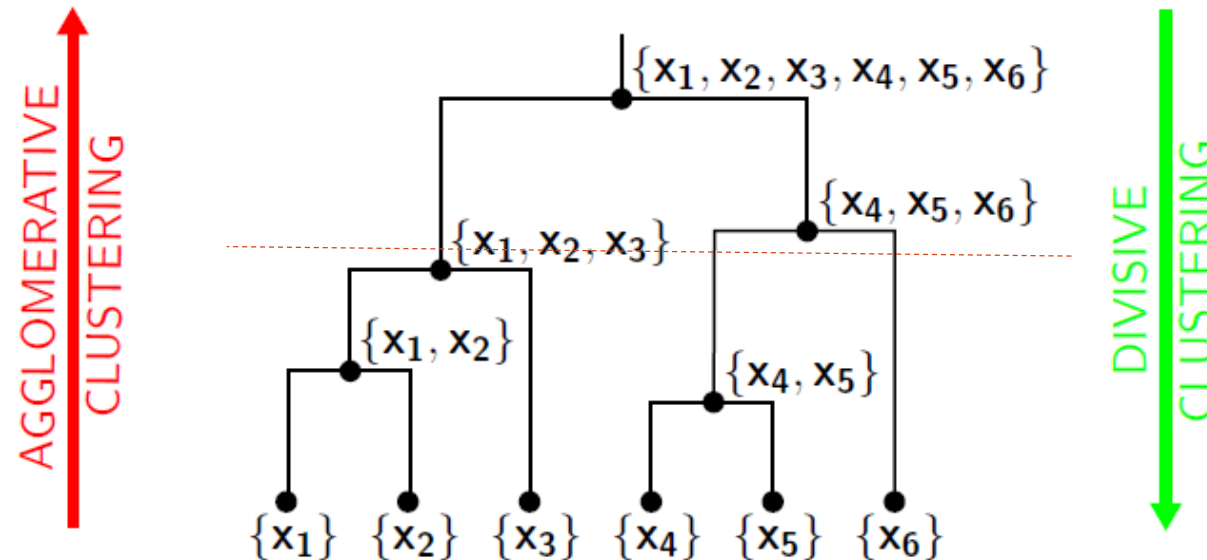


HIERARCHICAL AGGLOMERATIVE CLUSTERING (HAC) AND DENDROGRAM



CLUSTERING METHODS

- Graph methods
 - Single link
 - Complete link
 - Average
 - Weighted
- Geometric methods
 - Ward
 - Centroid
 - Median



DISSIMILARITY

- Some distance functions
 - Euclidean distance
 - Jaccard Similarity
 - Cosine Similarity
 - Edit distance
- Distance estimator
 - Model an estimator to be used with a predict probability function
 - Better for merging feature importances
- Dissimilarity (affinity) matrix




HIERARCHICAL CLUSTERING EXAMPLE (10F4)

We consider 5 data points in \mathbb{R}^2 :

- $x_1 = (1, 2)$
- $x_2 = (1, 2.5)$
- $x_3 = (3, 1)$
- $x_4 = (4, 0.5)$
- $x_5 = (4, 2)$

We consider the euclidean distance between data points.




HIERARCHICAL CLUSTERING EXAMPLE (20F4)

- Single Link method

```
import scipy.cluster.hierarchy as hac  
def __init__(self, method="single", affinity="euclidean")  
self.linkage_ = hac.linkage(X,  
                             method=self.method,  
                             metric=self.affinity)
```

$$D_{sl}(C_k, C_{k'}) = \min_{x \in C_k, y \in C_{k'}} \{D(x, y)\}$$



HIERARCHICAL CLUSTERING EXAMPLE (30F4)

$$D = D_{\text{eucl}} = D_{\text{sl}} = \begin{matrix} & \begin{matrix} x_1 & x_2 & x_3 & x_4 & x_5 \end{matrix} \\ \begin{matrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \end{matrix} & \begin{pmatrix} 0 & 0.5 & 2.24 & 3.35 & 3 \\ 0.5 & 0 & 2.5 & 3.61 & 3.04 \\ 2.24 & 2.5 & 0 & 1.12 & 1.41 \\ 3.35 & 3.61 & 1.12 & 0 & 1.5 \\ 3 & 3.04 & 1.41 & 1.5 & 0 \end{pmatrix} \end{matrix}$$

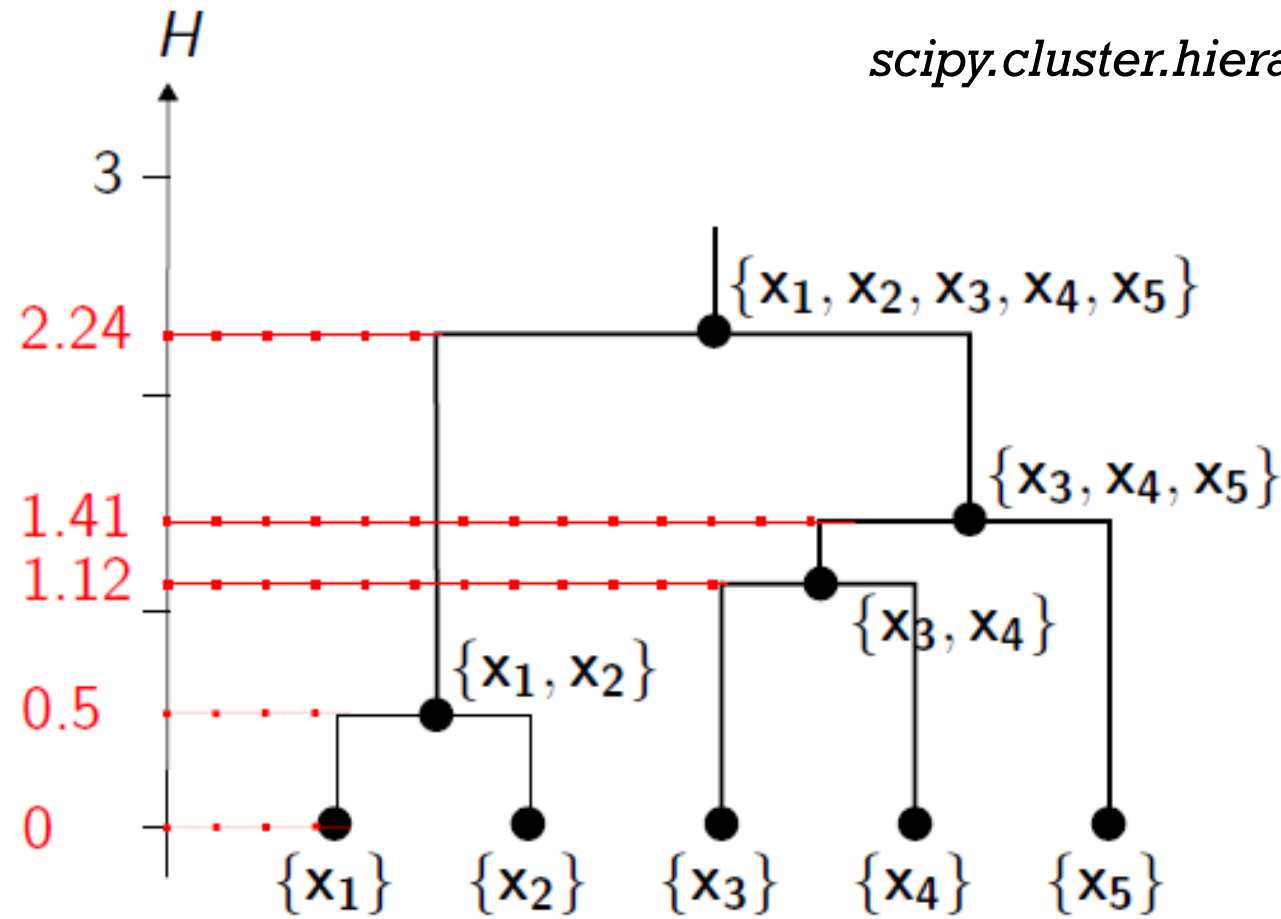
$$\text{dend}(h) = \{ \{x_1\}, \{x_2\}, \{x_3\}, \{x_4\}, \{x_5\} \quad \text{if } 0 \leq h$$

Merge x_1 and x_2

$$\text{dend}(h) = \begin{cases} \{x_1\}, \{x_2\}, \{x_3\}, \{x_4\}, \{x_5\} & \text{if } 0 \leq h < 0.5 \\ \{x_1, x_2\}, \{x_3\}, \{x_4\}, \{x_5\} & \text{if } 0.5 \leq h \end{cases}$$



HIERARCHICAL CLUSTERING EXAMPLE (40F4)





NUMBER OF CLUSTERS

- Flat cutting forming partitioned clustering from HAC
- Semi-supervised clustering
- DBSCAN

labels = hac.fcluster(self.linkage_, threshold)




NUMBER OF CLUSTERS (SEMI-SUPERVISED)




1 and 5 must be in the same cluster $\rightarrow n_{\text{clust}} = 2$





IDENTIFYING A MACHINE LEARNING PROBLEM



TASKS MAPPED TO ML (1 OF 2)

- Classification, recognition and completion
 - Decision trees
 - Supervised Neural Networks
 - Graph models: Hidden Markov Model, Bayesian Nets
- Optimization
 - Genetic Algorithms
 - Numerical Optimization and regression
- Smart Control
 - Fuzzy rules



TASKS MAPPED TO ML (2 OF 2)

- Clustering, grouping and segmentation
 - Unsupervised Neural Networks
 - HAC, K-Means, DBSCAN, Graph-based clustering
- Simulation
 - Multi-agent systems
- Relations
 - Graph networks
 - Collaborative filtering
 - Association rules and frequent sets



SOME ML APPLICATIONS (1 OF 2)

- Text mining: Topic and concept extraction → Unsupervised: Clustering, Graph and ontologies. Semi-supervised
- Pattern recognition: Character, speech and image recognition → Supervised: Classification
- Cross selling and personalized recommendation → Collaborative filtering, Association rules, frequent sets
- Churn management and customer retention → classification and clustering
- Games → search for optimal solution and multi-agent systems



SOME ML APPLICATIONS (2 OF 2)

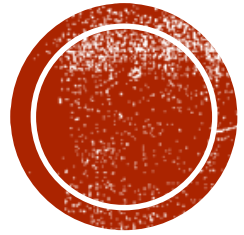
- Financials: risk management in insurance fraud detection → Graph, classification, clustering
- Security: Intrusion detection, spam filtering → Supervised: Classification
- Social networks: friends suggestion → Graph
- Bioinformatics (DNA)
- High energy physics (Higgs)
- Computer vision, Robotics ...



UNDERSTANDING THE DATA

- General statistics
- Correlations
- Plotting
- Understanding the source and the business behind the data
- Well-understanding of the requirements and the goal of the project





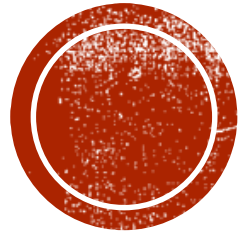
DATA PREPARATION



DATA CHALLENGES

- Feature extraction (What is a good feature?)
- Categories vs. continuous features vs. mixed
- Features Correlation
- Dimensionality reduction (the less the better for generalization)
- Missing values
- Normalization
- Outliers
- Feature selection
- Sampling
- Unbalanced dataset





RECOMMENDED PYTHON LIBRARIES



RECOMMENDED PYTHON LIBRARIES


- Scipy
- Numpy (vectorized data operations)
- Matplotlib (plotting)
- Scikit-Learn
- Pandas (data frames)
- NetworkX (Graphs analysis)
- Pylearn2 (deep learning and neural networks)
- BEARD (entity recognition) – *to be presented within the use case*



SCIKIT-LEARN API

- `fit`
- *`predict`*
- *`fit_transform`*
- *`Transformers`*
- *`Pipelines`*
- *`Features union`*
- *`Feature importances`*
- *`@property`*
- Joblib and parallel computing





CASE STUDY: AUTHOR NAME DISAMBIGUATION

CERN Digital Library

ENTITY RESOLUTION, DUPLICATION

Real World



Digital World



Source: <http://www.datacommunitydc.org/blog/2013/08/entity-resolution-for-big-data>



WHAT IS THE PROBLEM?

Please meet Yang Yang, Yang Yang, and Yang Yang!

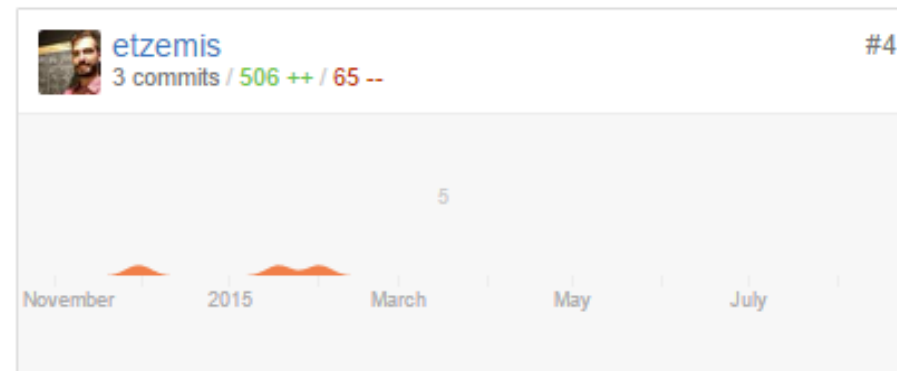
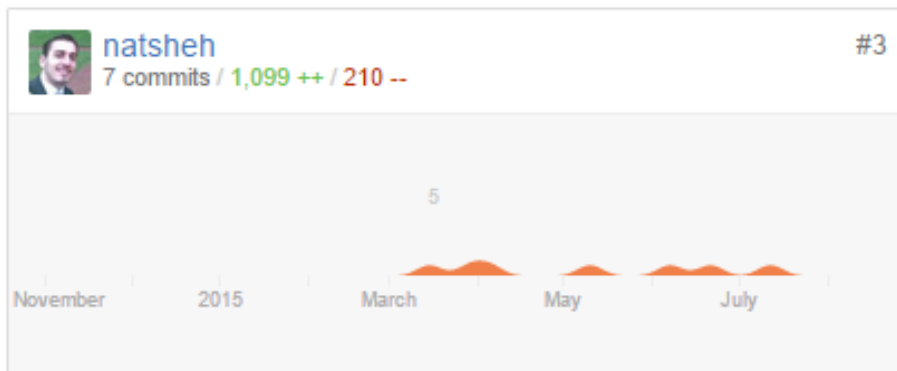
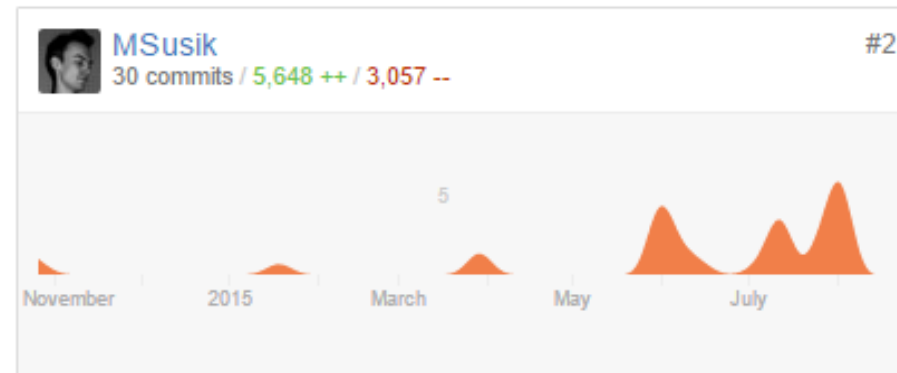


Source: Laura Rueda Garcia, CERN




OPEN SOURCE PROJECT @ CERN: BEARD

Beard is a Python library of machine learning tools for Bibliographic Entity Automatic Recognition and Disambiguation.

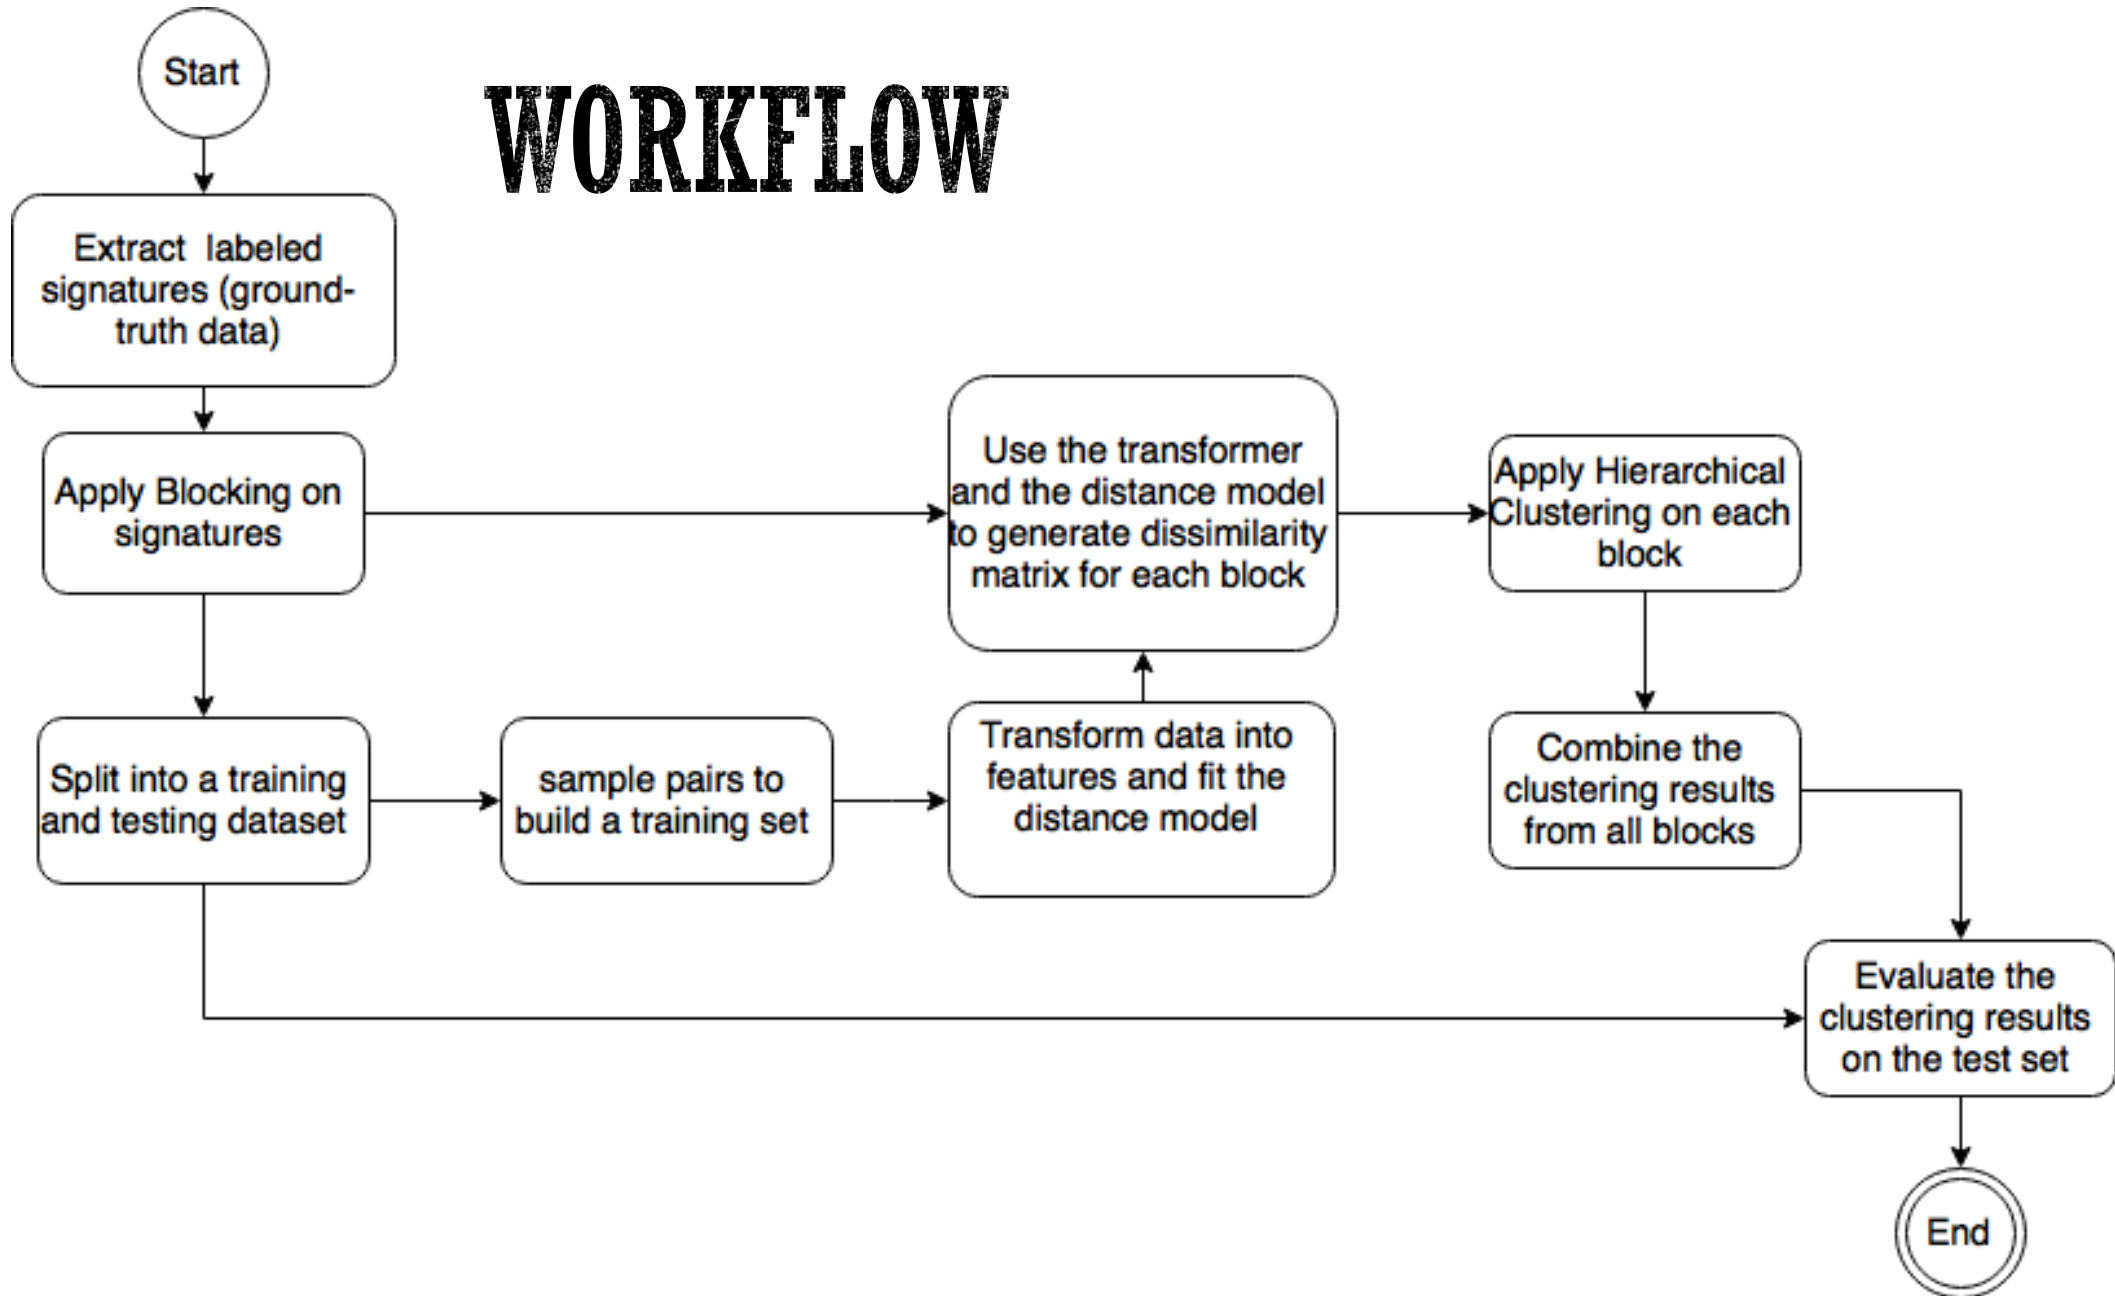


DATA OVERVIEW

- Signature: Unique occurrence of author name and publication id
- We have over:
 - 1.07 million publications
 - 9.27 millions signatures
 - 1.2 million claimed signatures (Ground-truth)
- Pair of signature (sig1, sig2, {0,1})
- Need for deviding into blocks



WORKFLOW



TRANSFORMER AND FEATURE UNION

```
def _build_distance_estimator(X, y, verbose=0, ethnicity_estimator=None):
    """Build a vector representation of a pair of signatures."""
    transformer = FeatureUnion([
        ("author_full_name_similarity", Pipeline([
            ("pairs", PairTransformer(element_transformer=Pipeline([
                ("full_name", FuncTransformer(func=get_author_full_name)),
                ("shaper", Shaper(newshape=(-1,))),
                ("tf-idf", TfidfVectorizer(analyzer="char_wb",
                                         ngram_range=(2, 4),
                                         dtype=np.float32,
                                         decode_error="replace")),
            ]), groupby=group_by_signature)),
            ("combiner", CosineSimilarity())
        ])),
        ("author_second_initial_similarity", Pipeline([
            ("pairs", PairTransformer(element_transformer=FuncTransformer(
                func=get_second_initial
            ), groupby=group_by_signature)),
            ("combiner", StringDistance(
                similarity_function="character_equality"))
        ])),
    ]),
```



FEATURE COULD BE ESTIMATED

```
))),  
("year_diff", Pipeline([  
    ("pairs", FuncTransformer(func=get_year, dtype=np.int)),  
    ("combiner", AbsoluteDifference())  
]))]
```

```
if ethnicity_estimator is not None:  
    transformer.transformer_list.append(("author_ethnicity", Pipeline([  
        ("pairs", PairTransformer(element_transformer=Pipeline([  
            ("name", FuncTransformer(func=get_author_full_name)),  
            ("shaper", Shaper(newshape=(-1,))),  
            ("classifier", EstimatorTransformer(ethnicity_estimator)),  
        ]), groupby=group_by_signature)),  
        ("sigmoid", FuncTransformer(func=expit)),  
        ("combiner", ElementMultiplication())  
    ]))
```



ETHNICITY ESTIMATOR: EXTERNAL DATA

```
# Load data
data = pd.read_csv(args.input_datafile)
y = data.RACE.values
X = ["%s, %s" % (last, first) for last, first in zip(data.NAMELAST.values,
                                                    data.NAMEFIRST.values)]
X = [normalize_name(name) for name in X]

# Train an estimator
estimator = Pipeline([
    ("transformer", TfidfVectorizer(analyzer="char_wb",
                                   ngram_range=(1, 5),
                                   min_df=0.00005,
                                   dtype=np.float32,
                                   decode_error="replace")),
    ("classifier", LinearSVC(C=args.C))])
estimator.fit(X, y)
```



PIPELINE OF TRANSFORMER AND CLASSIFIER

```
# Train a classifier on these vectors
classifier = GradientBoostingClassifier(n_estimators=500,
                                       max_depth=9,
                                       max_features=10,
                                       learning_rate=0.125,
                                       verbose=verbose)
```

```
# Return the whole pipeline
estimator = Pipeline([("transformer", transformer),
                      ("classifier", classifier)]).fit(X, y)
```

```
return estimator
```

```
def learn_model(distance_pairs, input_signatures, input_records,
               distance_model, verbose=0, ethnicity_estimator=None):
    """Learn the distance model for pairs of signatures.
```



BLOCK CLUSTERING

```
clusterer = BlockClustering(  
    blocking=block_last_name_first_initial,  
    base_estimator=ScipyHierarchicalClustering(  
        affinity=_affinity,  
        threshold=clustering_threshold,  
        method=clustering_method,  
        supervised_scoring=b3_f_score),  
    verbose=verbose,  
    n_jobs=n_jobs).fit(X, y)  
  
labels = clusterer.labels_
```



HAC CLUSTERING

```
import scipy.cluster.hierarchy as hac

from sklearn.base import BaseEstimator
from sklearn.base import ClusterMixin

class ScipyHierarchicalClustering(BaseEstimator, ClusterMixin):

    """Wrapper for Scipy's hierarchical clustering implementation.

    Attributes
    -----
    labels_ : ndarray, shape (n_samples,)
        Array of labels assigned to the input data.

    linkage_ : ndarray
        The linkage matrix.
    """

    def __init__(self, method="single", affinity="euclidean",
                 threshold=None, n_clusters=None, criterion="distance",
                 depth=2, R=None, monocrit=None, unsupervised_scoring=None,
                 supervised_scoring=None, scoring_data=None):
```



PROPERTY

```
@property
def labels_(self):
    """Compute the labels assigned to the input data.

    Note that labels are computed on-the-fly from the linkage matrix,
    based on the value of self.threshold or self.n_clusters.
    """
    n_clusters = self.n_clusters


    if n_clusters is not None:
        if n_clusters < 1 or n_clusters > self.n_samples_:
            raise ValueError("n_clusters must be within [1; n_samples].")
        else:
            thresholds = np.concatenate(([0],
                                         self.linkage[:, 2],
                                         [self.linkage[-1, 2]]))

            for i in range(len(thresholds) - 1):
                t1, t2 = thresholds[i:i + 2]
                threshold = (t1 + t2) / 2.0
                labels = hac.fcluster(self.linkage_, threshold,
                                     criterion=self.criterion,
                                     depth=self.depth, R=self.R,
                                     monocrit=self.monocrit)

                if len(np.unique(labels)) == n_clusters:
                    _, labels = np.unique(labels, return_inverse=True)
                    return labels
```

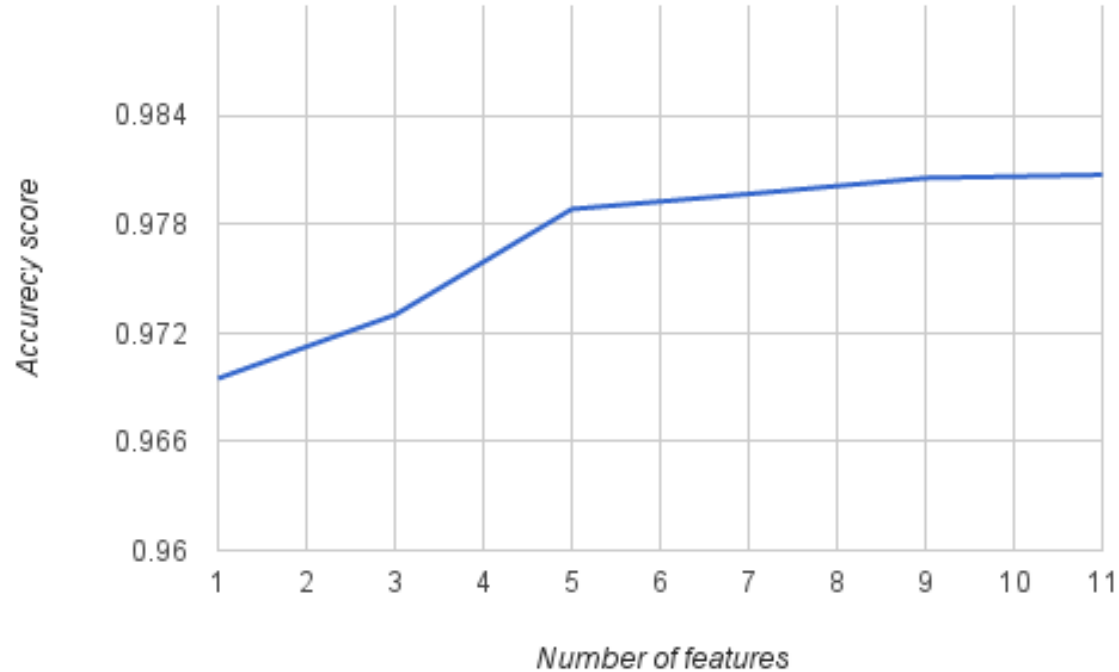


FEATURE IMPORTANCES

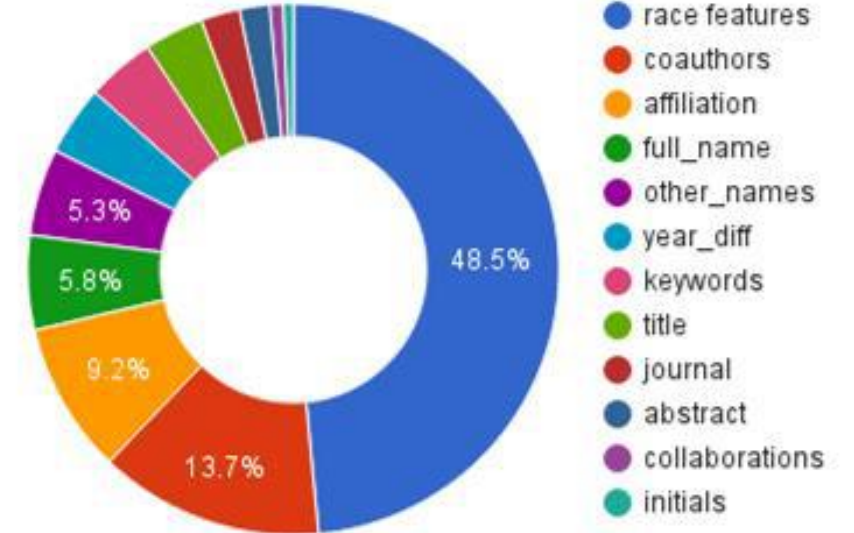


RECURSIVE REMOVAL


Recursive Removal of Least Important Features

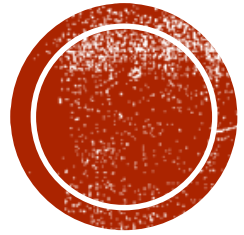


Feature / Importances



CLUSTERING METHODS





REFERENCES AND RECOMMENDED READINGS



REFERENCES

- Data Clustering- Part 2, M2 DMKM Slides, Julien Ah-Pine, Université Lyon 2, 2014
- BEARD: <https://github.com/inveniosoftware/beard>
- Scikit-learn: <http://scikit-learn.org>
- Hussein, AL-NATSHEH. "Bibliographic Entity Automatic Recognition and Disambiguation.", CERN, 2015
<https://preprints.cern.ch/record/2036112/files/CERN-THESIS-2015-098.pdf>
- Buitinck, Lars, Gilles Louppe, Mathieu Blondel, Fabian Pedregosa, Andreas Mueller, Olivier Grisel, Vlad Niculae et al. "API design for machine learning software: experiences from the scikit-learn project." *arXiv preprint arXiv:1309.0238* (2013).
<http://arxiv.org/pdf/1309.0238.pdf>



RECOMMENDED ML COURSES



Stanford University
Machine Learning



University of Washington
Introduction to Data Science

Practice on:

kaggle



CONTACT DETAILS:

- <https://ch.linkedin.com/in/natsheh>
- <https://twitter.com/hnatsheh>
- h.natsheh@ciapple.com



HUSSEIN AL-NATSEH

Thanks for attending!

